

HUMAN EMOTION DETECTION USING NUERAL NETWORK

A PROJECT BY
GAURAV KUMAR SINGH

Abstract

Human emotion recognition plays an important role in the relational relationship. The automatic recognition of emotions has been a functioning examination theme from early times. Hence, there are a few advances made in this field. Emotions are reflected from discourse, hand and signals of the body and through outward appearances. Henceforth extracting and understanding of emotion has a high significance of the collaboration among human and machine correspondence. This project depicts the advances made in this field and the different methodologies utilized for recognition of emotions. The main objective of the paper is to propose a real-time implementation of emotion recognition system.

Contents

| | |
|--|----|
| 1. Introduction..... | 01 |
| 2. Problem Statement..... | 01 |
| 3. Aim of the Project..... | 01 |
| 4. Literature Survey..... | 02 |
| 4.1. Dataset Used..... | 02 |
| 4.2. Attribute Details..... | 02 |
| 4.3. Existing System..... | 04 |
| 5. Design and Implementation..... | 04 |
| 5.1. Proposed System..... | 04 |
| 5.2. Training Phase..... | 06 |
| 5.3. Classification and Detection..... | 10 |
| 5.3.1. Single Face Detection..... | 10 |
| 5.3.2. Multiple Face Detection..... | 12 |
| 6. Code... .. | 13 |
| 7. Conclusion..... | 21 |
| 8. References..... | 21 |

1. Introduction

One of the current top applications of artificial intelligence using neural networks is the recognition of faces in photos and videos. Most techniques process visual data and search for general patterns present in human faces. Face recognition can be used for surveillance purposes by law enforcers as well as in crowd management. Other present-day applications involve automatic blurring of faces on Google Street view footage and automatic recognition of Facebook friends in photos. An even more advanced development in this field is emotion recognition. In addition to only identifying faces, the computer uses the arrangement and shape of e.g. eyebrows and lips to determine the facial expression and hence the emotion of a person. One possible application for this lies in the area of surveillance and behavioural analysis by law enforcement. The success of service robotics decisively depends on a smooth robot to user interaction. Thus, a robot should be able to extract information just from the face of its user, e.g. identify the emotional state or deduce gender. Interpreting correctly any of these elements using machine learning (ML) techniques has proven to be complicated due the high variability of the samples within each task. This leads to models with millions of parameters trained under thousands of samples. Furthermore, the human accuracy for classifying an image of a face in one of 7 different emotions is $65\% \pm 5\%$. One can observe the difficulty of this task by trying to manually classify the FER-2013 dataset images in Figure 1 within the following classes {"angry", "disgust", "fear", "happy", "sad", "surprise", "neutral"}.

2. Problem Statement

Many actions can be defined via observing the emotion the person is expressing. One of the current top applications of artificial intelligence using neural networks is the recognition of faces in photos and videos. Most techniques process visual data and search for general patterns present in human faces.

3. Aim of the Project

Face recognition can be used for surveillance purposes by law enforcers as well as in crowd management. Other present-day applications involve automatic blurring of faces on Google Street view footage and automatic recognition of Facebook friends in photos. An even more advanced development in this field is emotion recognition. In addition to only identifying faces, the computer uses the arrangement and shape of e.g. eyebrows and lips to determine the facial expression and hence the emotion of a person. One possible application for this lies in the area of surveillance and behavioural analysis by law enforcement.

4. Literature Survey

4.1. Dataset Used

The dataset has been taken from Kaggle on the URL: <https://www.kaggle.com/deadskull7/fer2013>. The data consists of 48x48 pixel grayscale images of faces. The faces have been automatically registered so that the face is more or less cantered and occupies about the same amount of space in each image. The task is to categorize each face based on the emotion shown in the facial expression in to one of seven categories (0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral). The training set consists of 28,709 examples. The public test set used for the leader board consists of 3,589 examples. The final test set, which was used to determine the winner of the competition, consists of another 3,589 examples. This dataset was prepared by Pierre-Luc Carrier and Aaron Courville, as part of an ongoing research project. They have graciously provided the workshop organizers with a preliminary version of their dataset to use for this contest.



Fig. 1: Samples of the FER-2013 emotion dataset.

4.2. Attribute Details

The dataset has total 3 attributes out of which 2 are independent variables and one is the dependent variable i.e. target variable which classifies the image into an emotion. There are 2 independent variables:

1. **Pixels:** The 48x48 grey-scale image is made in an array of 48x48 entries where each pixel is represented by its intensity.
2. **Usage:** Categorized into 3 parts Training, Public testing and private testing
3. **emotion:** Containing no. 0 to 6 for the emotions: 0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral

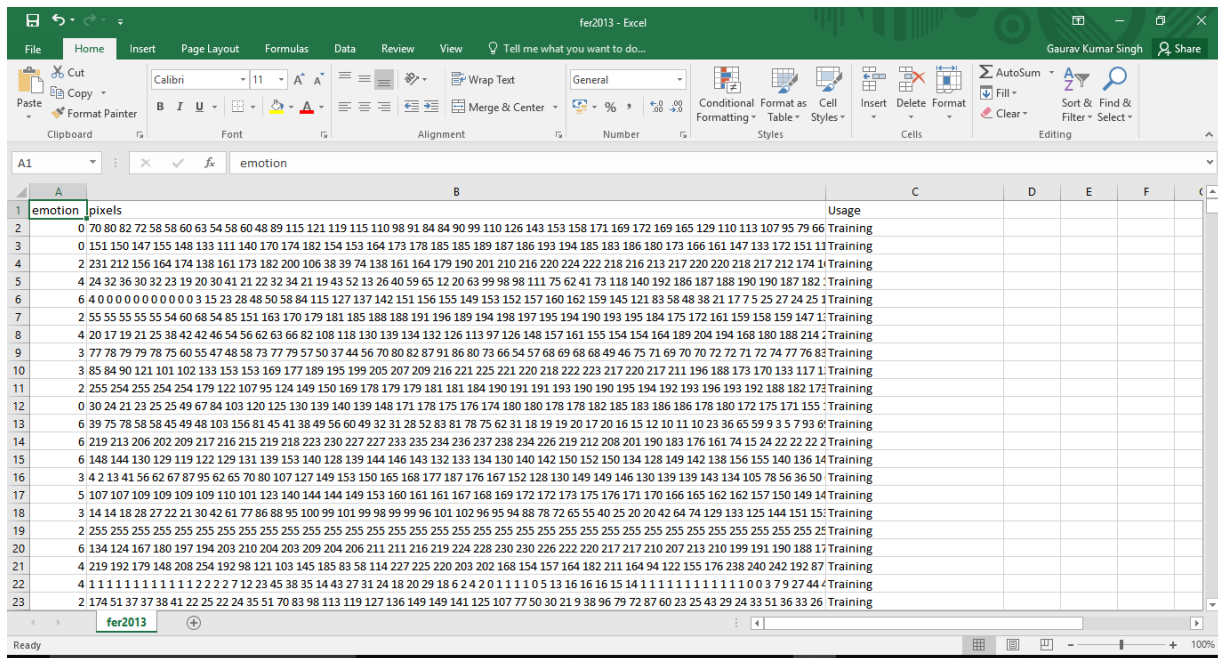


Fig. 2: Screenshot of the fer2-13 dataset.

4.3. Existing Systems

| Comparative Study | | | | |
|---|--|---|--|--|
| Title | Technique | Database | Performance (%) | Remarks |
| Statistical Moments based Facial expression Analysis | Feature Extraction: Zernike moments Classification: Naive Bayesian classifier | JAFFE (Japanese Female Facial expression) database 60 images used for experiment. | Average accuracy for six emotions is 81.66% in time less than 2 seconds. | Emotion accuracy graph shows highest recognition rate of happiness and lowest recognition rate of sadness. |
| Facial expression recognition with Auto-Illumination correction | Expressions on the face are determined with Action Units (AU's) | Single and Multiple face image | 60% recognition rate for multiple face image | Illumination on image plays vital role. |
| Identification-driven Emotion recognition system for a Social Robot | Hybrid approach used for personalized emotion recognition, | MUG facial expression database used. More than 50 people frontal face database used aged between 20-25 years. | 82% performance achieved with KNN Classifiers. | 3D model facial image used.KNN classifier gives good performance for emotion recognition. |
| The application study of learner's face detection and location in the teaching network system based on emotion | SVM(Support Vector Machine) classifier based Adaboost algorithm used | PIE face image database used | Detection and Correction rate 95% or more. | Presents application of face emotion recognition with of E-learning system. |

5. Design and Implementation

5.1. Proposed System

A Convolutional neural networks with pooling layers is used to detect the several emotions. The model relies on the idea of eliminating completely the fully connected layers. The architecture combines the deletion of the fully connected layer and the inclusion of the

combined depth-wise separable convolutions and residual modules. I have also implement multiple face emotion recognition which can detect multiple faces on a single frame. Major steps involved are:

- First, the use of small CNNs alleviate us from slow performances in hardware-constrained systems such robot platforms.
- And second, the reduction of parameters provides a better generalization under an Occam's razor framework.
- The model relies on the idea of eliminating completely the fully connected layers.
- The architecture combines the deletion of the fully connected layer and the inclusion of the combined depth-wise separable convolutions and residual modules.
- Both architectures were trained with the ADAM optimizer use Batch Normalization for better accuracy.
- Convolutional 2D network has been used to train the models.
- Take the real-time video or video file as an input
- Put it under processing
- Output is displayed.

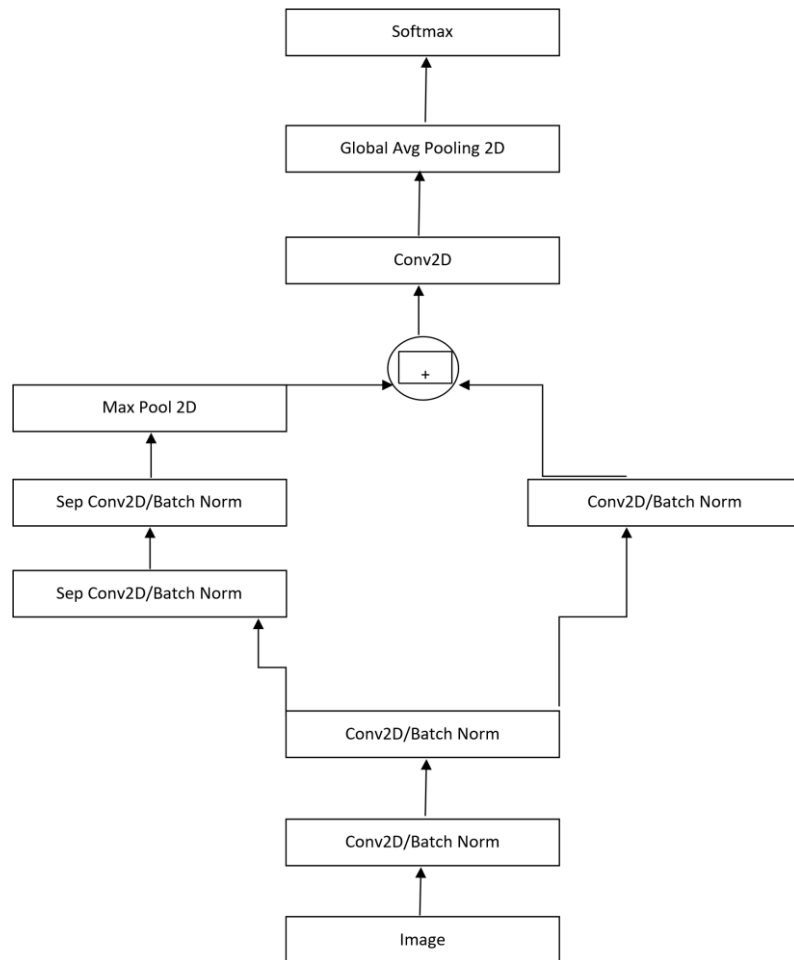


Fig. 3: Architecture of the proposed system.

5.2. Training Phase

The model was trained with the fer2013 dataset and the model achieves an accuracy of 66%. The training process ran for 106 epochs in which it garnered the accuracy. A new model was made only when the accuracy increased and error decreased. The following are the screenshots of the training process:

The screenshot shows the Spyder Python IDE with the file `train_emotion_classifier.py` open. The script defines a Keras model for emotion classification. The Python console displays the model summary, showing the first few layers:

| Layer (type) | Output Shape | Param # | Connected to |
|---|--------------------|---------|-----------------------------|
| input_1 (InputLayer) | (None, 48, 48, 1) | 0 | |
| conv2d_1 (Conv2D) | (None, 46, 46, 8) | 72 | input_1[0][0] |
| batch_normalization_1 (BatchNormalizer) | (None, 46, 46, 8) | 32 | conv2d_1[0][0] |
| activation_1 (Activation) | (None, 46, 46, 8) | 0 | batch_normalization_1[0][0] |
| conv2d_2 (Conv2D) | (None, 44, 44, 8) | 576 | activation_1[0][0] |
| batch_normalization_2 (BatchNormalizer) | (None, 44, 44, 8) | 32 | conv2d_2[0][0] |
| activation_2 (Activation) | (None, 44, 44, 8) | 0 | batch_normalization_2[0][0] |
| separable_conv2d_1 (SeparableConv2D) | (None, 44, 44, 16) | 280 | activation_2[0][0] |
| batch_normalization_4 (BatchNormalizer) | (None, 44, 44, 16) | 64 | separable_conv2d_1[0][0] |
| activation_3 (Activation) | (None, 44, 44, 16) | 0 | |

The screenshot shows the same Spyder Python IDE with the file `train_emotion_classifier.py` open. The Python console displays the continuation of the model summary, showing the final layers and parameter counts:

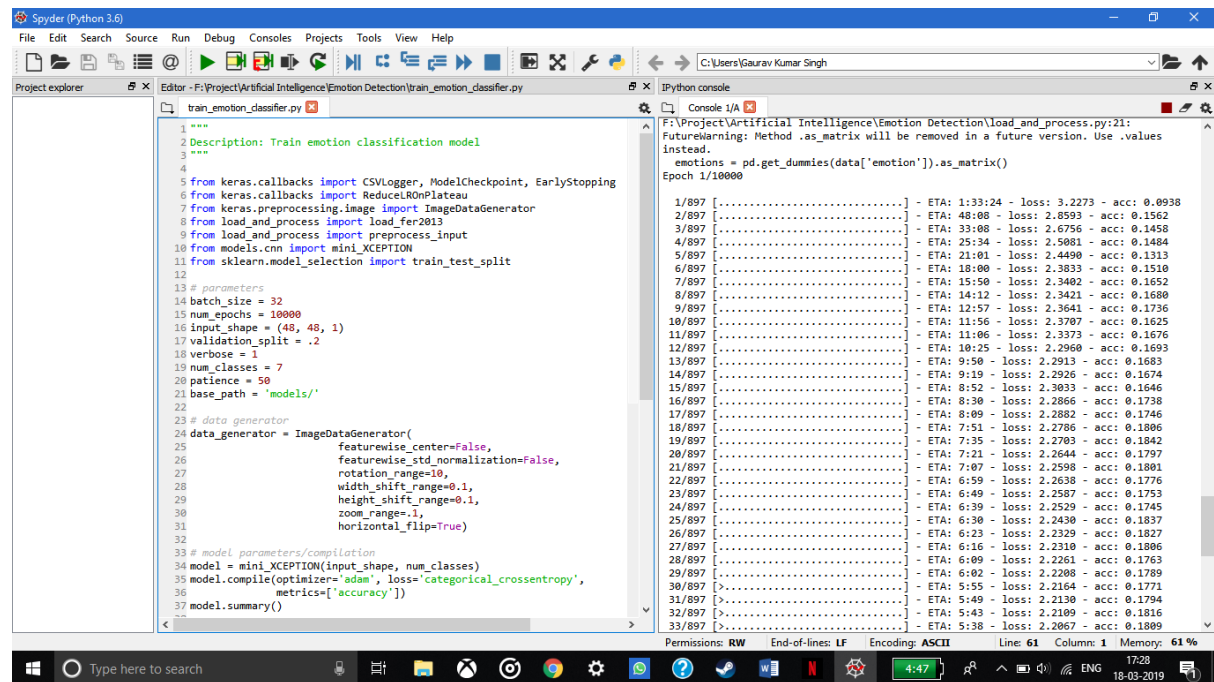
| Layer (type) | Output Shape | Param # | Connected to |
|---|-------------------|---------|----------------------------------|
| batch_normalization_14 (BatchNormalizer) | (None, 6, 6, 128) | 512 | separable_conv2d_8[0][0] |
| conv2d_6 (Conv2D) | (None, 3, 3, 128) | 8192 | batch_normalization_14[0][0] |
| max_pooling2d_4 (MaxPooling2D) | (None, 3, 3, 128) | 0 | conv2d_6[0][0] |
| batch_normalization_12 (BatchNormalizer) | (None, 3, 3, 128) | 512 | max_pooling2d_4[0][0] |
| add_4 (Add) | (None, 3, 3, 128) | 0 | batch_normalization_12[0][0] |
| conv2d_7 (Conv2D) | (None, 3, 3, 7) | 8071 | add_4[0][0] |
| global_average_pooling2d_1 (GlobalAveragePooling2D) | (None, 7) | 0 | conv2d_7[0][0] |
| predictions (Activation) | (None, 7) | 0 | global_average_pooling2d_1[0][0] |

Summary statistics:

- Total params: 58,423
- Trainable params: 56,951
- Non-trainable params: 1,472

The above screenshots describe the parameters of the training phase.

The following are the actual training process starting with epoch 1:



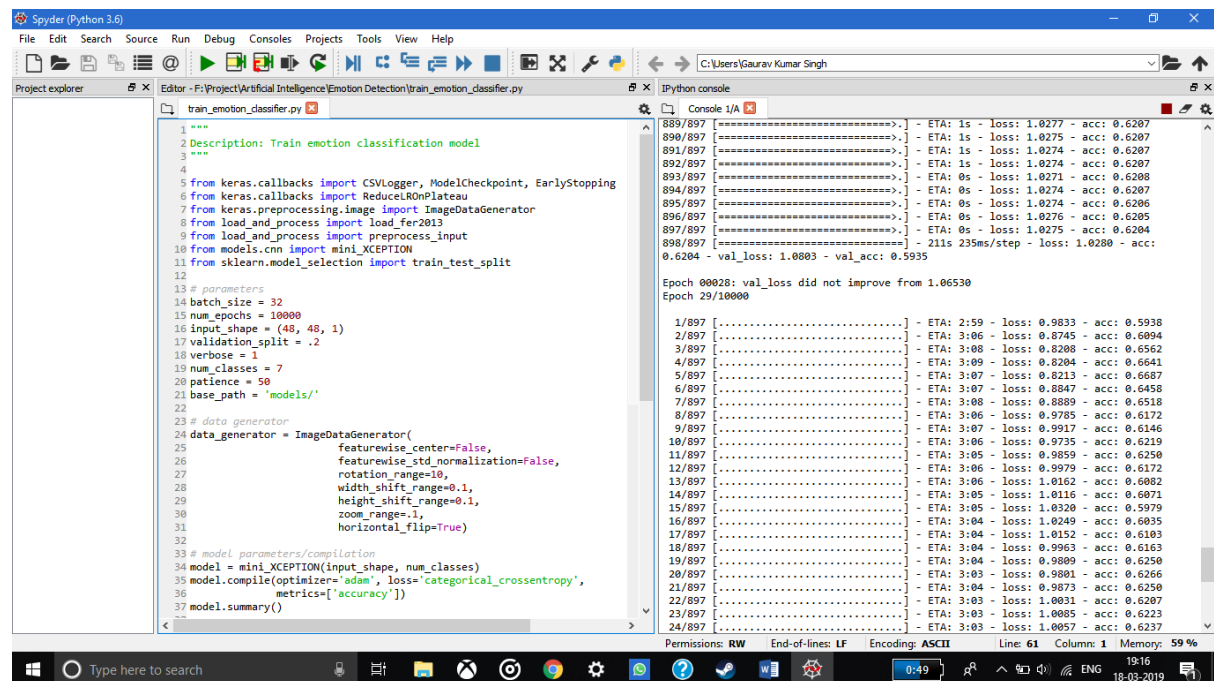
```
1 """
2 Description: Train emotion classification model
3 """
4
5 from keras.callbacks import CSVLogger, ModelCheckpoint, EarlyStopping
6 from keras.callbacks import ReduceLRonPlateau
7 from keras.preprocessing.image import ImageDataGenerator
8 from load_and_process import load_fer2013
9 from load_and_process import preprocess_input
10 from models.cnn import mini_XCEPTION
11 from sklearn.model_selection import train_test_split
12
13 # parameters
14 batch_size = 32
15 num_epochs = 10000
16 input_shape = (48, 48, 1)
17 validation_split = .2
18 verbose = 1
19 num_classes = 7
20 patience = 50
21 base_path = 'models/'
22
23 # data generator
24 data_generator = ImageDataGenerator(
25     featurewise_center=False,
26     featurewise_std_normalization=False,
27     rotation_range=10,
28     width_shift_range=0.1,
29     height_shift_range=0.1,
30     zoom_range=1,
31     horizontal_flip=True)
32
33 # model parameters/compilation
34 model = mini_XCEPTION(input_shape, num_classes)
35 model.compile(optimizer='adam', loss='categorical_crossentropy',
36               metrics=['accuracy'])
37 model.summary()
```

FutureWarning: Method .as_matrix will be removed in a future version. Use .values instead.

emotions = pd.get_dummies(data['emotion']).as_matrix()

Epoch 1/10000

| Batch | ETA | loss | acc |
|--------|---------|--------|--------|
| 1/897 | 1:33:24 | 3.2273 | 0.0938 |
| 2/897 | 48:08 | 2.8593 | 0.1562 |
| 3/897 | 33:08 | 2.6756 | 0.1458 |
| 4/897 | 25:34 | 2.5081 | 0.1484 |
| 5/897 | 21:01 | 2.4400 | 0.1313 |
| 6/897 | 18:00 | 2.3833 | 0.1510 |
| 7/897 | 15:50 | 2.3402 | 0.1652 |
| 8/897 | 14:12 | 2.3421 | 0.1680 |
| 9/897 | 12:57 | 2.3641 | 0.1736 |
| 10/897 | 11:56 | 2.3707 | 0.1625 |
| 11/897 | 11:06 | 2.3373 | 0.1676 |
| 12/897 | 10:25 | 2.2960 | 0.1693 |
| 13/897 | 9:50 | 2.2913 | 0.1683 |
| 14/897 | 9:19 | 2.2926 | 0.1674 |
| 15/897 | 8:52 | 2.3033 | 0.1646 |
| 16/897 | 8:30 | 2.2866 | 0.1738 |
| 17/897 | 8:09 | 2.2882 | 0.1746 |
| 18/897 | 7:51 | 2.2786 | 0.1806 |
| 19/897 | 7:35 | 2.2703 | 0.1842 |
| 20/897 | 7:21 | 2.2644 | 0.1797 |
| 21/897 | 7:07 | 2.2598 | 0.1801 |
| 22/897 | 6:59 | 2.2638 | 0.1776 |
| 23/897 | 6:49 | 2.2587 | 0.1753 |
| 24/897 | 6:39 | 2.2529 | 0.1745 |
| 25/897 | 6:30 | 2.2430 | 0.1837 |
| 26/897 | 6:23 | 2.2329 | 0.1827 |
| 27/897 | 6:16 | 2.2310 | 0.1806 |
| 28/897 | 6:09 | 2.2261 | 0.1763 |
| 29/897 | 6:02 | 2.2208 | 0.1789 |
| 30/897 | 5:55 | 2.2164 | 0.1771 |
| 31/897 | 5:49 | 2.2130 | 0.1794 |
| 32/897 | 5:43 | 2.2109 | 0.1816 |
| 33/897 | 5:38 | 2.2067 | 0.1809 |



```
1 """
2 Description: Train emotion classification model
3 """
4
5 from keras.callbacks import CSVLogger, ModelCheckpoint, EarlyStopping
6 from keras.callbacks import ReduceLRonPlateau
7 from keras.preprocessing.image import ImageDataGenerator
8 from load_and_process import load_fer2013
9 from load_and_process import preprocess_input
10 from models.cnn import mini_XCEPTION
11 from sklearn.model_selection import train_test_split
12
13 # parameters
14 batch_size = 32
15 num_epochs = 10000
16 input_shape = (48, 48, 1)
17 validation_split = .2
18 verbose = 1
19 num_classes = 7
20 patience = 50
21 base_path = 'models/'
22
23 # data generator
24 data_generator = ImageDataGenerator(
25     featurewise_center=False,
26     featurewise_std_normalization=False,
27     rotation_range=10,
28     width_shift_range=0.1,
29     height_shift_range=0.1,
30     zoom_range=1,
31     horizontal_flip=True)
32
33 # model parameters/compilation
34 model = mini_XCEPTION(input_shape, num_classes)
35 model.compile(optimizer='adam', loss='categorical_crossentropy',
36               metrics=['accuracy'])
37 model.summary()
```

Epoch 00028: val_loss did not improve from 1.06530

Epoch 29/10000

| Batch | ETA | loss | acc |
|--------|------|--------|--------|
| 1/897 | 2:59 | 0.9833 | 0.5938 |
| 2/897 | 3:06 | 0.8745 | 0.6094 |
| 3/897 | 3:08 | 0.8208 | 0.6562 |
| 4/897 | 3:09 | 0.8204 | 0.6641 |
| 5/897 | 3:07 | 0.8213 | 0.6687 |
| 6/897 | 3:07 | 0.8847 | 0.6458 |
| 7/897 | 3:08 | 0.8889 | 0.6518 |
| 8/897 | 3:06 | 0.9785 | 0.6172 |
| 9/897 | 3:07 | 0.9917 | 0.6146 |
| 10/897 | 3:06 | 0.9735 | 0.6219 |
| 11/897 | 3:05 | 0.9859 | 0.6250 |
| 12/897 | 3:06 | 0.9979 | 0.6172 |
| 13/897 | 3:06 | 1.0162 | 0.6082 |
| 14/897 | 3:05 | 1.0116 | 0.6071 |
| 15/897 | 3:05 | 1.0320 | 0.5979 |
| 16/897 | 3:04 | 1.0249 | 0.6035 |
| 17/897 | 3:04 | 1.0152 | 0.6103 |
| 18/897 | 3:04 | 0.9963 | 0.6163 |
| 19/897 | 3:04 | 0.9809 | 0.6250 |
| 20/897 | 3:03 | 0.9801 | 0.6266 |
| 21/897 | 3:04 | 0.9873 | 0.6250 |
| 22/897 | 3:03 | 1.0031 | 0.6207 |
| 23/897 | 3:03 | 1.0085 | 0.6223 |
| 24/897 | 3:03 | 1.0057 | 0.6237 |

In the above figure the val_loss doesn't improve hence no model is made.

```
1 """
2 Description: Train emotion classification model
3 """
4
5 from keras.callbacks import CSVLogger, ModelCheckpoint, EarlyStopping
6 from keras.callbacks import ReduceLRonPlateau
7 from keras.preprocessing.image import ImageDataGenerator
8 from load_and_process import load_fer2013
9 from load_and_process import preprocess_input
10 from models.cnn import mini_XCEPTION
11 from sklearn.model_selection import train_test_split
12
13 # parameters
14 batch_size = 32
15 num_epochs = 10000
16 input_shape = (48, 48, 1)
17 validation_split = .2
18 verbose = 1
19 num_classes = 7
20 patience = 50
21 base_path = 'models/'
22
23 # data generator
24 data_generator = ImageDataGenerator(
25     featurewise_center=False,
26     featurewise_std_normalization=False,
27     rotation_range=10,
28     width_shift_range=0.1,
29     height_shift_range=0.1,
30     zoom_range=.1,
31     horizontal_flip=True)
32
33 # model parameters/compilation
34 model = mini_XCEPTION(input_shape, num_classes)
35 model.compile(optimizer='adam', loss='categorical_crossentropy',
36             metrics=['accuracy'])
37 model.summary()
```

Epoch 00045: val_loss did not improve from 1.00387
Epoch 46/10000

| Epoch | ETA | loss | acc |
|---------|------|------------|---|
| 883/897 | 3s | 0.9827 | 0.6657 |
| 884/897 | 3s | 0.9827 | 0.6656 |
| 885/897 | 2s | 0.9829 | 0.6656 |
| 886/897 | 2s | 0.9828 | 0.6656 |
| 887/897 | 2s | 0.9826 | 0.6658 |
| 888/897 | 2s | 0.9825 | 0.6658 |
| 889/897 | 1s | 0.9825 | 0.6658 |
| 890/897 | 1s | 0.9825 | 0.6658 |
| 891/897 | 1s | 0.9830 | 0.6657 |
| 892/897 | 1s | 0.9832 | 0.6655 |
| 893/897 | 0s | 0.9833 | 0.6655 |
| 894/897 | 0s | 0.9829 | 0.6657 |
| 895/897 | 0s | 0.9833 | 0.6655 |
| 896/897 | 0s | 0.9832 | 0.6656 |
| 897/897 | 0s | 0.9834 | 0.6656 |
| 898/897 | 235s | 262ms/step | loss: 0.9828 - acc: 0.6657 - val_loss: 1.0045 - val_acc: 0.6329 |

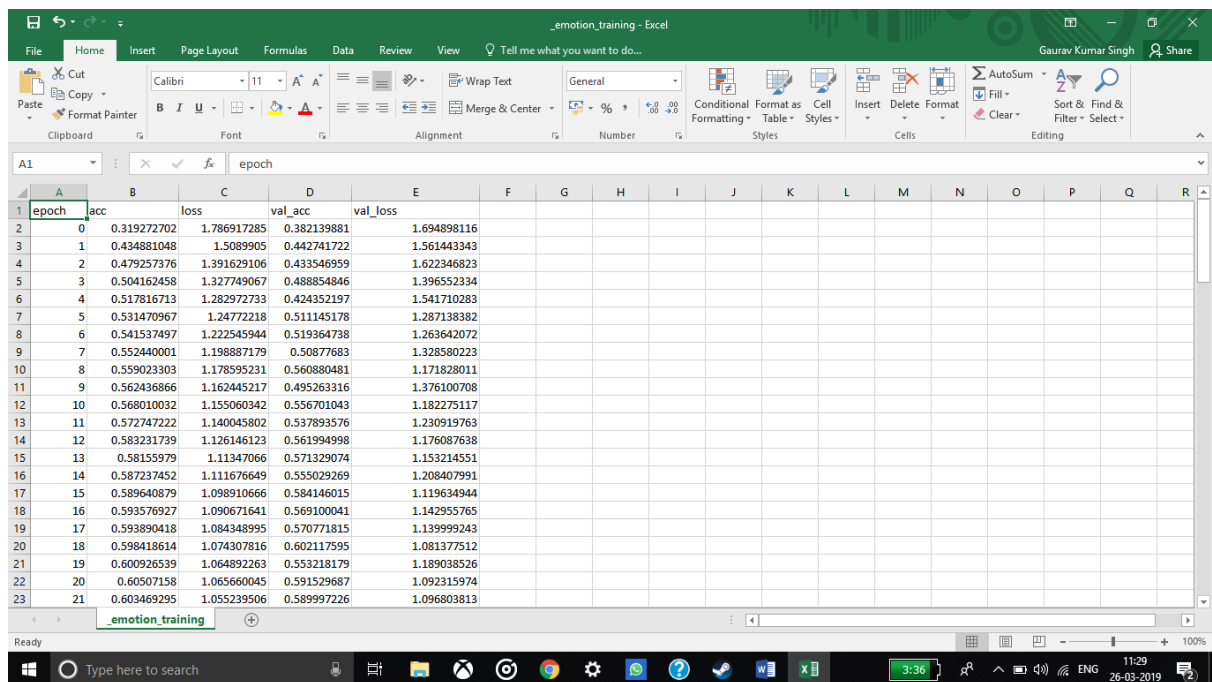
```
1 """
2 Description: Train emotion classification model
3 """
4
5 from keras.callbacks import CSVLogger, ModelCheckpoint, EarlyStopping
6 from keras.callbacks import ReduceLRonPlateau
7 from keras.preprocessing.image import ImageDataGenerator
8 from load_and_process import load_fer2013
9 from load_and_process import preprocess_input
10 from models.cnn import mini_XCEPTION
11 from sklearn.model_selection import train_test_split
12
13 # parameters
14 batch_size = 32
15 num_epochs = 10000
16 input_shape = (48, 48, 1)
17 validation_split = .2
18 verbose = 1
19 num_classes = 7
20 patience = 50
21 base_path = 'models/'
22
23 # data generator
24 data_generator = ImageDataGenerator(
25     featurewise_center=False,
26     featurewise_std_normalization=False,
27     rotation_range=10,
28     width_shift_range=0.1,
29     height_shift_range=0.1,
30     zoom_range=.1,
31     horizontal_flip=True)
32
33 # model parameters/compilation
34 model = mini_XCEPTION(input_shape, num_classes)
35 model.compile(optimizer='adam', loss='categorical_crossentropy',
36             metrics=['accuracy'])
37 model.summary()
```

0.6752 - val_loss: 1.0010 - val_acc: 0.6328
Epoch 00063: val_loss improved from 1.00228 to 1.00099, saving model to models/_mini_XCEPTION.63-0.63.hdf5
Epoch 64/10000

| Epoch | ETA | loss | acc |
|--------|------|--------|--------|
| 1/897 | 3:13 | 0.9792 | 0.6562 |
| 2/897 | 3:20 | 0.8115 | 0.7344 |
| 3/897 | 3:19 | 0.9801 | 0.6979 |
| 4/897 | 3:21 | 0.9887 | 0.6562 |
| 5/897 | 3:19 | 0.9536 | 0.6750 |
| 6/897 | 3:20 | 0.9724 | 0.6510 |
| 7/897 | 3:20 | 0.9248 | 0.6607 |
| 8/897 | 3:18 | 0.9147 | 0.6602 |
| 9/897 | 3:19 | 0.9311 | 0.6493 |
| 10/897 | 3:18 | 0.9249 | 0.6562 |
| 11/897 | 3:18 | 0.9878 | 0.6648 |
| 12/897 | 3:17 | 0.9888 | 0.6693 |
| 13/897 | 3:17 | 0.9166 | 0.6635 |
| 14/897 | 3:17 | 0.9135 | 0.6652 |
| 15/897 | 3:17 | 0.9275 | 0.6684 |
| 16/897 | 3:17 | 0.9217 | 0.6660 |
| 17/897 | 3:16 | 0.9110 | 0.6654 |
| 18/897 | 3:16 | 0.9113 | 0.6667 |

Here in the above figure the model is made as the val_loss improves.

The training log file maintained during the training phase:



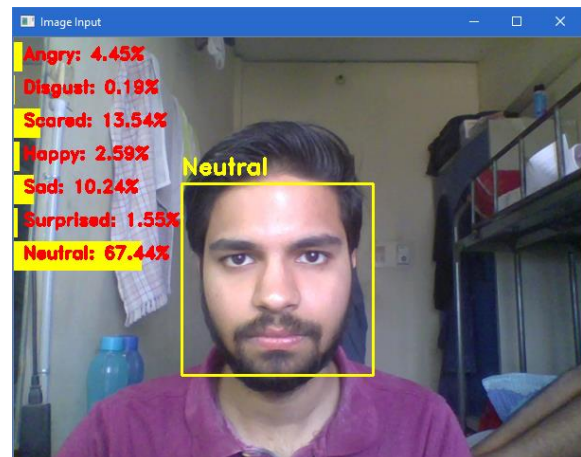

| epoch | acc | loss | val_acc | val_loss |
|-------|-------------|-------------|-------------|-------------|
| 0 | 0.319272702 | 1.786917285 | 0.382139881 | 1.694898116 |
| 1 | 0.434881048 | 1.5089905 | 0.442741722 | 1.561443343 |
| 2 | 0.479257376 | 1.391629106 | 0.433546959 | 1.622346823 |
| 3 | 0.504162458 | 1.327749067 | 0.488854846 | 1.396552334 |
| 4 | 0.517816713 | 1.282972733 | 0.424352197 | 1.541710283 |
| 5 | 0.531470967 | 1.24772218 | 0.511145178 | 1.287138382 |
| 6 | 0.541537497 | 1.222545944 | 0.519364738 | 1.263642072 |
| 7 | 0.552440001 | 1.198887179 | 0.50877683 | 1.328580223 |
| 8 | 0.559203303 | 1.178595231 | 0.560880481 | 1.171828011 |
| 9 | 0.562436866 | 1.162445217 | 0.495263316 | 1.376100708 |
| 10 | 0.568010032 | 1.155060342 | 0.556701043 | 1.182275117 |
| 11 | 0.572747222 | 1.140045802 | 0.537893576 | 1.230919763 |
| 12 | 0.583231739 | 1.126146123 | 0.561994998 | 1.176087638 |
| 13 | 0.58155979 | 1.11347066 | 0.571329074 | 1.153214551 |
| 14 | 0.587237452 | 1.111676649 | 0.555029269 | 1.208407991 |
| 15 | 0.589640879 | 1.098910666 | 0.584146015 | 1.119634944 |
| 16 | 0.593576927 | 1.090671641 | 0.569100041 | 1.142955765 |
| 17 | 0.593890418 | 1.084348995 | 0.570771815 | 1.139999243 |
| 18 | 0.598418614 | 1.074307816 | 0.602117595 | 1.081377512 |
| 19 | 0.600926539 | 1.064892263 | 0.553218179 | 1.189038526 |
| 20 | 0.60507158 | 1.065660045 | 0.591529687 | 1.092315974 |
| 21 | 0.603469295 | 1.055239506 | 0.589997226 | 1.096803813 |


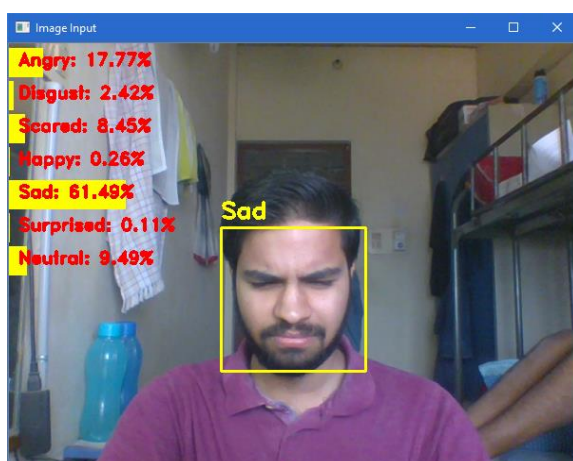
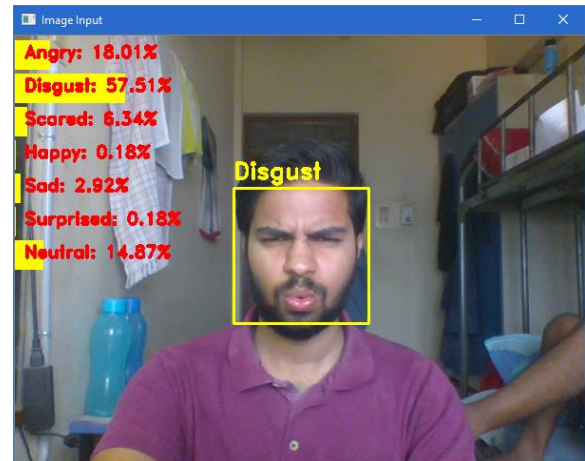
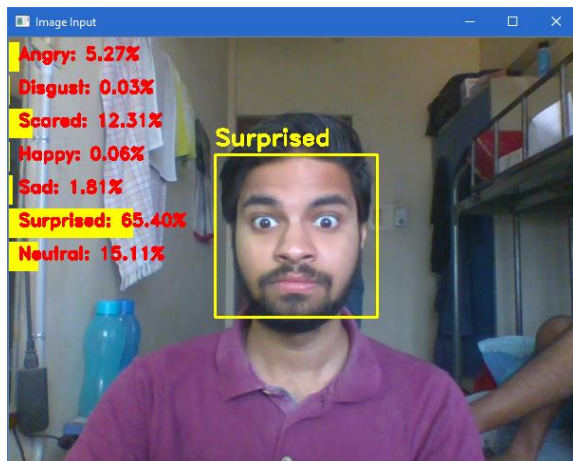

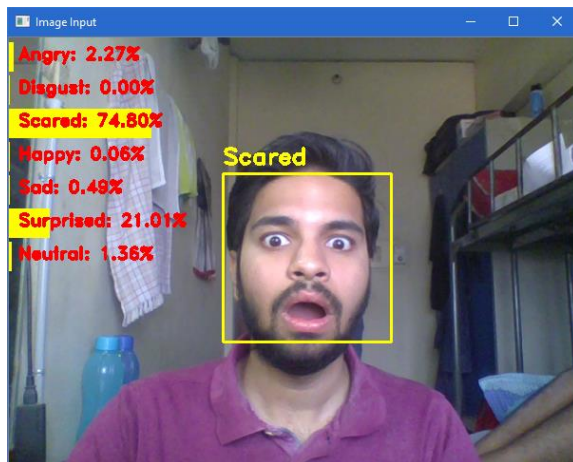
5.2. Classification and Detection

After the model is created we take the input and predict the emotion using the model. There are two types of programs. The first one detects a single face and also draws the probabilities of different emotion in the same frame as a function of the probabilities in a bar graph form. The second one detects and classifies multiple face simultaneously without displaying probabilities.

5.2.1. Single Face Detection

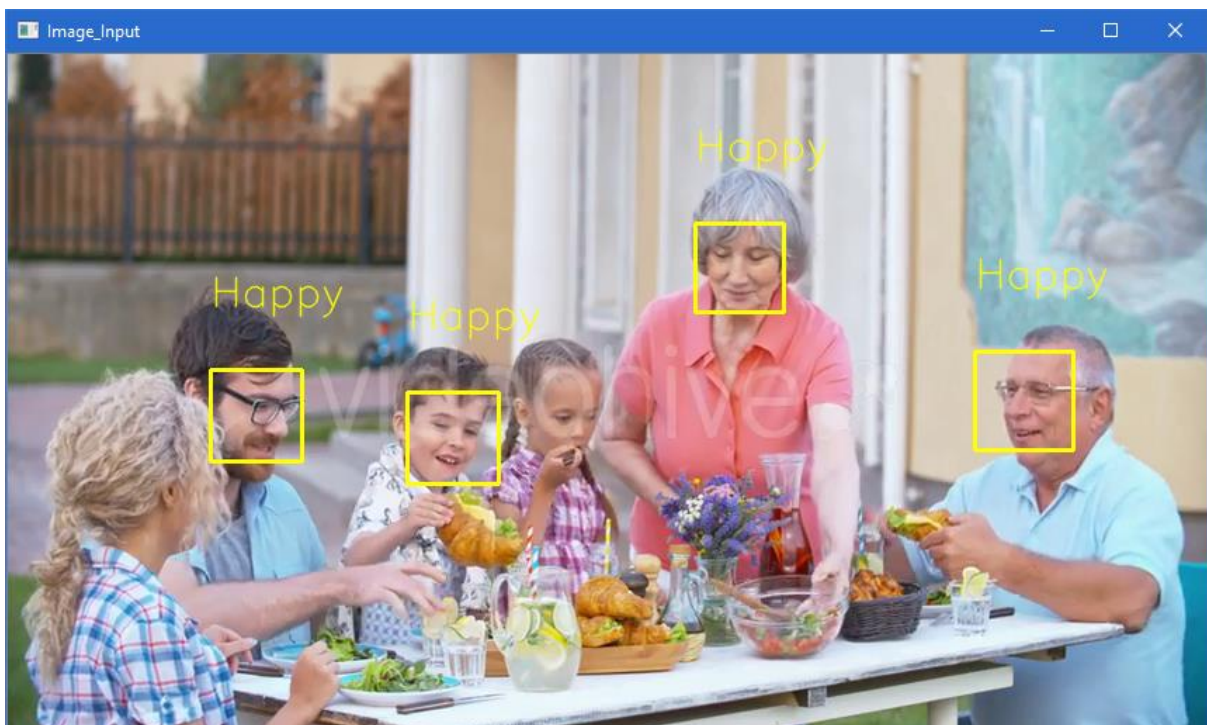

The following are the screenshots of the single face detection:





5.2.2. Multiple Face Detection

The following are the screenshots of the multiple face detection:



6. Code

load_and_process:

```
import pandas as pd

import cv2

import numpy as np


dataset_path = 'Dataset/fer2013/fer2013.csv'

image_size=(48,48)


def load_fer2013():

    data = pd.read_csv(dataset_path)

    pixels = data['pixels'].tolist()

    width, height = 48, 48

    faces = []

    for pixel_sequence in pixels:

        face = [int(pixel) for pixel in pixel_sequence.split(' ')]

        face = np.asarray(face).reshape(width, height)

        face = cv2.resize(face.astype('uint8'),image_size)

        faces.append(face.astype('float32'))

    faces = np.asarray(faces)

    faces = np.expand_dims(faces, -1)

    emotions = pd.get_dummies(data['emotion']).as_matrix()

    return faces, emotions


def preprocess_input(x, v2=True):

    x = x.astype('float32')

    x = x / 255.0

    if v2:

        x = x - 0.5

        x = x * 2.0
```

```
return x
```

train_emotion_classifier:

```
"""
```

Description: Train emotion classification model

```
"""
```

```
from keras.callbacks import CSVLogger, ModelCheckpoint, EarlyStopping
```

```
from keras.callbacks import ReduceLROnPlateau
```

```
from keras.preprocessing.image import ImageDataGenerator
```

```
from load_and_process import load_fer2013
```

```
from load_and_process import preprocess_input
```

```
from models.cnn import mini_XCEPTION
```

```
from sklearn.model_selection import train_test_split
```

```
# parameters
```

```
batch_size = 32
```

```
num_epochs = 10000
```

```
input_shape = (48, 48, 1)
```

```
validation_split = .2
```

```
verbose = 1
```

```
num_classes = 7
```

```
patience = 50
```

```
base_path = 'models/'
```

```
# data generator
```

```
data_generator = ImageDataGenerator(
```

```
    featurewise_center=False,
```

```
    featurewise_std_normalization=False,
```

```
    rotation_range=10,
```

```
    width_shift_range=0.1,
```



```

        height_shift_range=0.1,

        zoom_range=.1,

        horizontal_flip=True)

# model parameters/compilation

model = mini_XCEPTION(input_shape, num_classes)

model.compile(optimizer='adam', loss='categorical_crossentropy',

               metrics=['accuracy'])

model.summary()


# callbacks

log_file_path = base_path + '_emotion_training.log'

csv_logger = CSVLogger(log_file_path, append=False)

early_stop = EarlyStopping('val_loss', patience=patience)

reduce_lr = ReduceLROnPlateau('val_loss', factor=0.1,

                               patience=int(patience/4), verbose=1)

trained_models_path = base_path + '_mini_XCEPTION'

model_names = trained_models_path + '.{epoch:02d}-{val_acc:.2f}.hdf5'

model_checkpoint = ModelCheckpoint(model_names, 'val_loss', verbose=1,

                                   save_best_only=True)

callbacks = [model_checkpoint, csv_logger, early_stop, reduce_lr]


# loading dataset

faces, emotions = load_fer2013()

faces = preprocess_input(faces)

num_samples, num_classes = emotions.shape

xtrain, xtest, ytrain, ytest = train_test_split(faces, emotions, test_size=0.2, shuffle=True)

model.fit_generator(data_generator.flow(xtrain, ytrain,

                                       batch_size),

                   steps_per_epoch=len(xtrain) / batch_size,

                   epochs=num_epochs, verbose=1, callbacks=callbacks,

                   validation_data=(xtest, ytest))

```

Single:

```
import cv2

from keras.preprocessing.image import img_to_array

import imutils

from keras.models import load_model

import numpy as np


# parameters for loading data and images

detection_model_path = 'haarcascade_files/haarcascade_frontalface_default.xml'

emotion_model_path = 'models/_mini_XCEPTION.102-0.66.hdf5'

# hyper-parameters for bounding boxes shape

# loading models

face_detection = cv2.CascadeClassifier(detection_model_path)

emotion_classifier = load_model(emotion_model_path, compile=False)

EMOTIONS = ["Angry" ,"Disgust","Scared" ,"Happy", "Sad", "Surprised", "Neutral"]


# starting video streaming

camera = cv2.VideoCapture(0)

while True:

    frame = camera.read()[1]

    #reading the frame

    frame = imutils.resize(frame,width=600)

    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)

    faces = face_detection.detectMultiScale(gray,scaleFactor=1.1,minNeighbors=5,minSize=(30,30),flags=cv2.CASCADE_SCALE_IMAGE)

    if len(faces) > 0:

        faces = sorted(faces, reverse=True,

            key=lambda x: (x[2] - x[0]) * (x[3] - x[1]))[0]
```

```

(fX, fY, fW, fH) = faces

    # Extract the ROI of the face from the grayscale image, resize it to a fixed 28x28 pixels, and then
prepare

    # the ROI for classification via the CNN

roi = gray[fY:fY + fH, fX:fX + fW]
roi = cv2.resize(roi, (64, 64))
roi = roi.astype("float") / 255.0
roi = img_to_array(roi)
roi = np.expand_dims(roi, axis=0)


preds = emotion_classifier.predict(roi)[0]
emotion_probability = np.max(preds)
label = EMOTIONS[preds.argmax()]


for (i, (emotion, prob)) in enumerate(zip(EMOTIONS, preds)):

    # construct the label text
    text = "{}: {:.2f}%".format(emotion, prob * 100)

    w = int(prob * 200)

    cv2.rectangle(frame, (0, (i * 35) + 5),
        (w, (i * 35) + 35), (0, 255, 255), -1)
    cv2.putText(frame, text, (10, (i * 35) + 23),
        cv2.FONT_HERSHEY_DUPLEX, 0.6,
        (0, 0, 255), 2)
    cv2.putText(frame, label, (fX, fY - 10),
        cv2.FONT_HERSHEY_DUPLEX, 0.8, (0, 255, 255), 2)
    cv2.rectangle(frame, (fX, fY), (fX + fW, fY + fH),
        (0, 255, 255), 2)

cv2.imshow('Image Input', frame)

```

```
if cv2.waitKey(1) & 0xFF == ord('q'):
    break
```

```
camera.release()
```

```
cv2.destroyAllWindows()
```

Multiple:

```
import cv2
```

```
import numpy as np
```

```
from keras.models import load_model
```

```
from statistics import mode
```

```
#from utils.datasets import get_labels
```

```
#from utils.inference import detect_faces
```

```
from utils.inference import draw_text
```

```
from utils.inference import draw_bounding_box
```

```
from utils.inference import apply_offsets
```

```
#from utils.inference import load_detection_model
```

```
from utils.preprocessor import preprocess_input
```

```
USE_WEBCAM = True
```

```
# If false, loads video file source
```

```
# parameters for loading data and images
```

```
emotion_model_path = 'models/_mini_XCEPTION.102-0.66.hdf5'
```

```
#emotion_model_path = 'models/_mini_XCEPTION.63-0.63.hdf5'
```

```
emotion_labels=["Angry" ,"Disgust","Scared" ,"Happy", "Sad", "Surprised", "Neutral"]
```

```
# hyper-parameters for bounding boxes shape
```

```
frame_window = 10
```

```
emotion_offsets = (20, 40)
```

```
# loading models

face_cascade = cv2.CascadeClassifier('./haarcascade_files/haarcascade_frontalface_default.xml')

emotion_classifier = load_model(emotion_model_path, compile=False)

# getting input model shapes for inference

emotion_target_size = emotion_classifier.input_shape[1:3]

# starting lists for calculating modes

emotion_window = []

# starting video streaming

cv2.namedWindow('Image_Input')

video_capture = cv2.VideoCapture(0)

# Select video or webcam feed

cap = None

if (USE_WEBCAM == True):

    cap = cv2.VideoCapture(0) # Webcam source

else:

    cap = cv2.VideoCapture('dinner.mp4') # Video file source

while cap.isOpened(): # True:

    ret, bgr_image = cap.read()

    #bgr_image = video_capture.read()[1]

    gray_image = cv2.cvtColor(bgr_image, cv2.COLOR_BGR2GRAY)

    rgb_image = cv2.cvtColor(bgr_image, cv2.COLOR_BGR2RGB)

    faces = face_cascade.detectMultiScale(gray_image, scaleFactor=1.1, minNeighbors=5,
```

```
minSize=(30, 30), flags=cv2.CASCADE_SCALE_IMAGE)
```

```
for face_coordinates in faces:
```

```
    x1, x2, y1, y2 = apply_offsets(face_coordinates, emotion_offsets)
```

```
    gray_face = gray_image[y1:y2, x1:x2]
```

```
    try:
```

```
        gray_face = cv2.resize(gray_face, (emotion_target_size))
```

```
    except:
```

```
        continue
```

```
    gray_face = preprocess_input(gray_face, True)
```

```
    gray_face = np.expand_dims(gray_face, 0)
```

```
    gray_face = np.expand_dims(gray_face, -1)
```

```
    emotion_prediction = emotion_classifier.predict(gray_face)
```

```
    emotion_probability = np.max(emotion_prediction)
```

```
    emotion_label_arg = np.argmax(emotion_prediction)
```

```
    emotion_text = emotion_labels[emotion_label_arg]
```

```
    emotion_window.append(emotion_text)
```

```
if len(emotion_window) > frame_window:
```

```
    emotion_window.pop(0)
```

```
try:
```

```
    emotion_mode = mode(emotion_window)
```

```
except:
```

```
    continue
```

```
color=np.asarray((255,255,0))
```

```
color = color.astype(int)
```

```
color = color.tolist()
```

```
draw_bounding_box(face_coordinates, rgb_image, color)
```

```
draw_text(face_coordinates, rgb_image, emotion_mode,
```

```
color, 0, -45, 1, 1)
```

```
bgr_image = cv2.cvtColor(rgb_image, cv2.COLOR_RGB2BGR)
```

```
cv2.imshow('Image_Input', bgr_image)
```

```
if cv2.waitKey(1) & 0xFF == ord('q'):
```

```
    break
```

```
cap.release()
```

```
cv2.destroyAllWindows()
```

9. Conclusion

A real time CNN model is built which can identify the emotion of the user standing in front of the camera. It can identify them via image or video input. The proposed architecture has been systematically built in order to reduce the amount of parameters. We began by eliminating completely the fully connected layers and by reducing the amount of parameters in the remaining convolutional layers via depth-wise separable convolutions.

10. References

- [1] Real-time Convolutional Neural Networks for Emotion and Gender Classification by Octavio Arriaga, Paul G. Ploger and Matias Valdenegro
- [2] Human Emotion Recognition System by Dilbag Singh
- [3] https://en.wikipedia.org/wiki/Emotion_recognition
- [4] https://docs.opencv.org/3.4.3/d7/d8b/tutorial_py_face_detection.html
- [5] <https://www.geeksforgeeks.org/image-classifier-using-cnn/>